

Speedy Q-Learning

Akif Khan

Reinforcement Learning Course Project
Indian Institute of Science

April, 2024



➊ Motivation for Speedy Q-Learning

➋ Algorithm

➌ Simulations

➍ Proposed Modification to SQL

➎ Results

➏ Questions

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

Q-Learning is Slow

- We know Q-Learning converges to the optimal value function with sample complexity $\tilde{O}\left(\frac{|S||A|}{(1-\gamma)^5\epsilon^2}\right)$

Q-Learning is Slow

- We know Q-Learning converges to the optimal value function with sample complexity $\tilde{O}\left(\frac{|S||A|}{(1-\gamma)^5\epsilon^2}\right)$
- But it is slow!

Q-Learning is Slow

- We know Q-Learning converges to the optimal value function with sample complexity $\tilde{O}\left(\frac{|S||A|}{(1-\gamma)^5\epsilon^2}\right)$
- But it is slow!
 - Sample inefficient due to decaying learning rate in the stochastic approximation. For ex. for $\alpha_k = \frac{1}{k+1}$, we have an overly pessimistic step size for large k

Q-Learning is Slow

- We know Q-Learning converges to the optimal value function with sample complexity $\tilde{O}\left(\frac{|S||A|}{(1-\gamma)^5\epsilon^2}\right)$
- But it is slow!
 - Sample inefficient due to decaying learning rate in the stochastic approximation. For ex. for $\alpha_k = \frac{1}{k+1}$, we have an overly pessimistic step size for large k
 - Propagation of the Bellman Operator throughout the space for discount factors close to 1

Q-Learning is Slow

- We know Q-Learning converges to the optimal value function with sample complexity $\tilde{O}\left(\frac{|S||A|}{(1-\gamma)^5\epsilon^2}\right)$
- But it is slow!
 - Sample inefficient due to decaying learning rate in the stochastic approximation. For ex. for $\alpha_k = \frac{1}{k+1}$, we have an overly pessimistic step size for large k
 - Propagation of the Bellman Operator throughout the space for discount factors close to 1
 - These issues can be resolved with Speedy Q-Learning

① Motivation for Speedy Q-Learning

② Algorithm

Notations

Pseudocode for Speedy Q-Learning

Why does Speedy Q-Learning Work?

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

① Motivation for Speedy Q-Learning

② Algorithm

Notations

Pseudocode for Speedy Q-Learning

Why does Speedy Q-Learning Work?

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

Notations Used

- X, A are finite sets and $Z = |X||A|$
- The immediate rewards $r(x, a)$ are bounded, i.e. $|r(x, a)| \leq R_{max}$
- Let γ be the discount factor and define $\beta = \frac{1}{1-\gamma}$ and $V_{max} = \beta R_{max}$
- T_k and T are the Bellman operator and the Bellman optimality operator
- $Q(x, a)$ is the Q value function and define $(MQ)(x) = \max_{a \in A} Q(x, a), \forall x \in X$
- α_k be the step size

① Motivation for Speedy Q-Learning

② Algorithm

Notations

Pseudocode for Speedy Q-Learning

Why does Speedy Q-Learning Work?

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

Pseudocode for Speedy Q-Learning

Algorithm 1: Synchronous Speedy Q-Learning (SQL)

Input: Initial action-value function Q_0 , discount factor γ , and number of iteration T

```

 $Q_{-1} := Q_0;$  // Initialization
for  $k := 0, 1, 2, 3, \dots, T - 1$  do // Main loop
     $\alpha_k := \frac{1}{k+1};$ 
    for each  $(x, a) \in \mathcal{Z}$  do
        Generate the next state sample  $y_k \sim P(\cdot | x, a);$ 
         $\mathcal{T}_k Q_{k-1}(x, a) := r(x, a) + \gamma M Q_{k-1}(y_k);$ 
         $\mathcal{T}_k Q_k(x, a) := r(x, a) + \gamma M Q_k(y_k);$  // Empirical Bellman operator
         $Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (\mathcal{T}_k Q_{k-1}(x, a) - Q_k(x, a)) + (1 - \alpha_k) (\mathcal{T}_k Q_k(x, a) - \mathcal{T}_k Q_{k-1}(x, a));$  // SQL update rule
    end
end
return  $Q_T$ 
  
```

Pseudocode for Speedy Q-Learning

Algorithm 1 Synchronous Speedy Q-Learning(SQL)

Initialize Q_0 , γ , Number of Iterations, T

$Q_{-1} \leftarrow Q_0$

for $k = 0, 1, \dots, T - 1$ **do**

$$\alpha_k = \frac{1}{k+1}$$

for each $(x, a) \in Z$ **do**

Generate the next sample $y_k \sim P(\cdot | x, a)$

$$T_k Q_{k-1}(x, a) \leftarrow r(x, a) + \gamma(MQ_{k-1})(y_k)$$

$$T_k Q_k(x, a) \leftarrow r(x, a) + \gamma(MQ_k)(y_k)$$

$$Q_{k+1}(x, a) \leftarrow Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) + (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

end

end

return Q_T

① Motivation for Speedy Q-Learning

② Algorithm

Notations

Pseudocode for Speedy Q-Learning

Why does Speedy Q-Learning Work?

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

An Intuitive Explanation

- The update rule for Q-Learning is given by

$$Q_{k+1}(x, a) = Q_k(x, a) + \alpha_k (T_k Q_k(x, a) - Q_k(x, a))$$

An Intuitive Explanation

- The update rule for Q-Learning is given by

$$Q_{k+1}(x, a) = Q_k(x, a) + \alpha_k (T_k Q_k(x, a) - Q_k(x, a))$$

- Adding and subtracting $T_k Q_{k-1}(x, a)$ for the term inside the bracket gives us

$$Q_{k+1}(x, a) = Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) + \alpha_k (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

An Intuitive Explanation

- The update rule for Q-Learning is given by

$$Q_{k+1}(x, a) = Q_k(x, a) + \alpha_k (T_k Q_k(x, a) - Q_k(x, a))$$

- Adding and subtracting $T_k Q_{k-1}(x, a)$ for the term inside the bracket gives us

$$Q_{k+1}(x, a) = Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) + \alpha_k (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

- The last term on the RHS goes to zero as $Q_k \rightarrow Q^*$, so a more aggressive step-size, $1 - \alpha_k = \frac{k}{k+1}$ can be used for this term

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL


⑤ Results

⑥ Questions

Environment and Libraries

Frozen-Lake environment:

- $4 \times 4 = 16$ states.
S:Start, H:Hole, G:Goal, F: Frozen
- 4 actions.
0:Left, 1:Down, 2: Right, 3:Up
- Rewards.
Reach Goal:+1, Reach Hole: 0, Reach Frozen: 0
- Slipping.
Slips from target direction to either side, with equal probabilities. Eg: If moving left, it goes to left with $\frac{1}{3}$ probability and slips top/bottom with probability $\frac{1}{3}$.



Libraries Used: gym, matplotlib, IPython.display and numpy.

Numerical Experiments for $\gamma = 0.99$

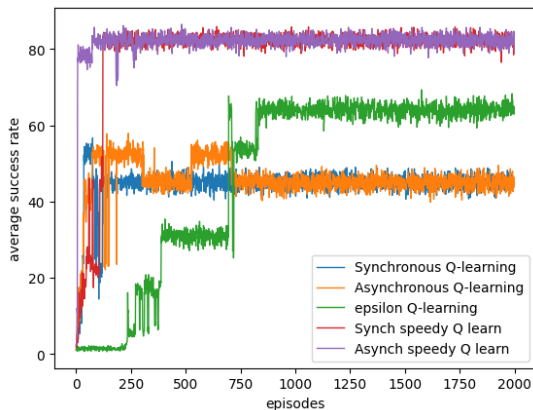


Figure 1: Discount factor $\gamma = 0.99$. Success Rate¹ vs Episodes²

¹Success is percentage of the trials reached to Goal.

²Number of Iterations

Numerical Experiments for $\gamma = 0.9$




Figure 2: Discount factor $\gamma = 0.9$. Success Rate³ vs Episodes⁴

³Success is percentage of the trials reached to Goal.

⁴Number of Iterations

Numerical Experiments for $\gamma = 0.8$




Figure 3: Discount factor $\gamma = 0.8$. Success Rate⁵ vs Episodes⁶

⁵Success is percentage of the trials reached to Goal.

⁶Number of Iterations

Numerical Experiments for $\gamma = 0.7$

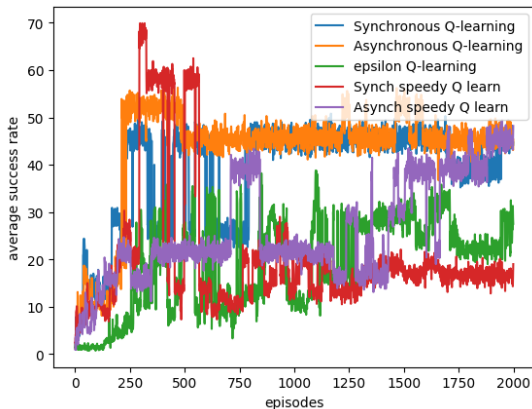


Figure 4: Discount factor $\gamma = 0.7$. Success Rate⁷ vs Episodes⁸

⁷Success is percentage of the trials reached to Goal.

⁸Number of Iterations

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

A Proposed Modification to SQL

- The Speedy Q-Learning update rule is

$$Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\ + (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

A Proposed Modification to SQL

- The Speedy Q-Learning update rule is

$$Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\ + (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

- Why to stop at one buffer, $Q_{k-1}(x, a)$? We can have more than one buffer like $Q_{k-2}(x, a)$, $Q_{k-3}(x, a)$ and so on. We can use a convex combination of α'_k 's such that the step-size gets more aggressive for the terms that go to zero faster

$$Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\ + \frac{2}{3} (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a)) \\ + \frac{1}{3} (1 - \alpha_k) (T_k Q_{k-1}(x, a) - T_k Q_{k-2}(x, a))$$

A Proposed Modification to SQL

- The Speedy Q-Learning update rule is

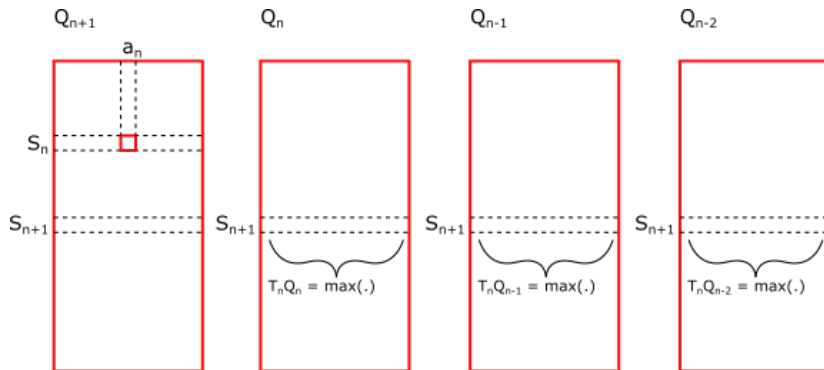
$$Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\ + (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a))$$

- Why to stop at one buffer, $Q_{k-1}(x, a)$? We can have more than one buffer like $Q_{k-2}(x, a)$, $Q_{k-3}(x, a)$ and so on. We can use a convex combination of α'_k 's such that the step-size gets more aggressive for the terms that go to zero faster

$$Q_{k+1}(x, a) := Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\ + \frac{2}{3} (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a)) \\ + \frac{1}{3} (1 - \alpha_k) (T_k Q_{k-1}(x, a) - T_k Q_{k-2}(x, a))$$

- Observation:** It seems to outperform the Speedy Q learning

A Proposed Modification to SQL



$$\begin{aligned}
 Q_{k+1}(x, a) &:= Q_k(x, a) + \alpha_k (T_k Q_{k-1}(x, a) - Q_k(x, a)) \\
 &\quad + \frac{2}{3} (1 - \alpha_k) (T_k Q_k(x, a) - T_k Q_{k-1}(x, a)) \\
 &\quad + \frac{1}{3} (1 - \alpha_k) (T_k Q_{k-1}(x, a) - T_k Q_{k-2}(x, a))
 \end{aligned}$$

Numerical results for Modified SQL(1 of 4)

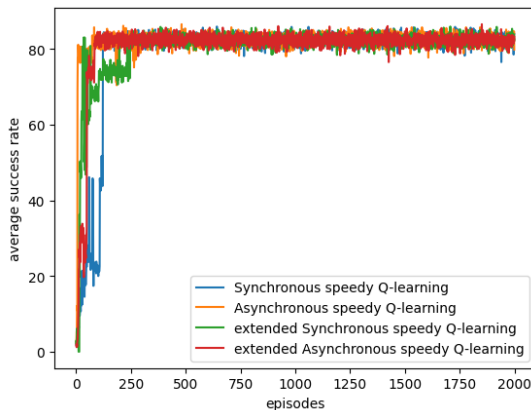


Figure 6: Discount factor $\gamma = 0.99$. Success Rate⁹ vs Episodes¹⁰

⁹Success is percentage of the trials reached to Goal.

¹⁰Number of Iterations

Numerical results for Modified SQL(2 of 4)

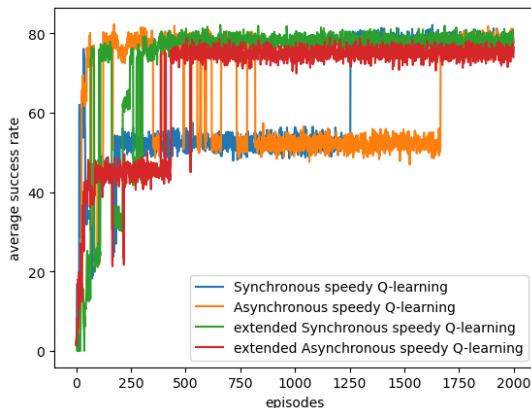


Figure 7: Discount factor $\gamma = 0.9$. Success Rate¹¹ vs Episodes¹²

¹¹Success is percentage of the trials reached to Goal.

¹²Number of Iterations

Numerical results for Modified SQL(3 of 4)




Figure 8: Discount factor $\gamma = 0.8$. Success Rate¹³ vs Episodes¹⁴

¹³Success is percentage of the trials reached to Goal.

¹⁴Number of Iterations

Numerical results for Modified SQL(4 of 4)

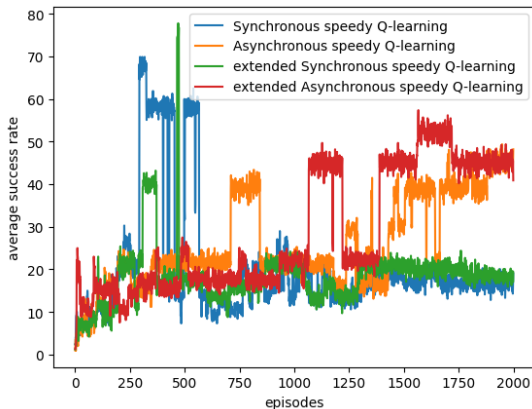


Figure 9: Discount factor $\gamma = 0.7$. Success Rate¹⁵ vs Episodes¹⁶

¹⁵Success is percentage of the trials reached to Goal.

¹⁶Number of Iterations

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

Main Theoretical Result

Comparison of Q-Learning vs Speedy Q-Learning

⑥ Questions

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

Main Theoretical Result

Comparison of Q-Learning vs Speedy Q-Learning

⑥ Questions

- Let Assumption 1 holds and T be a positive integer. Then, at iteration T of SQL with probability at least $1 - \delta$, we have

$$\|Q^* - Q_k\| \leq 2\beta^2 R_{\max} \left[\frac{\gamma}{T} + \sqrt{\frac{2 \log \frac{2n}{\delta}}{T}} \right]$$

- Let Assumption 1 holds and T be a positive integer. Then, at iteration T of SQL with probability at least $1 - \delta$, we have

$$\|Q^* - Q_k\| \leq 2\beta^2 R_{\max} \left[\frac{\gamma}{T} + \sqrt{\frac{2 \log \frac{2n}{\delta}}{T}} \right]$$

- This shows that $Q^* \xrightarrow{\text{a.s.}} Q$ with rate $\sqrt{\frac{1}{T}}$

- Let Assumption 1 holds and T be a positive integer. Then, at iteration T of SQL with probability at least $1 - \delta$, we have

$$\|Q^* - Q_k\| \leq 2\beta^2 R_{\max} \left[\frac{\gamma}{T} + \sqrt{\frac{2 \log \frac{2n}{\delta}}{T}} \right]$$

- This shows that $Q^* \xrightarrow{\text{a.s.}} Q$ with rate $\sqrt{\frac{1}{T}}$
- Additionally, for any $\epsilon > 0$, after $T = \frac{11.66\beta^4 R_{\max}^2 \log \frac{2n}{\delta}}{\epsilon^2}$ steps of SQL, $\|Q^* - Q_k\| \leq \epsilon$, with probability at least $1 - \delta$

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

Main Theoretical Result

Comparison of Q-Learning vs Speedy Q-Learning

⑥ Questions

- It has been shown for Q-Learning the finite time PAC bound for a step size, $\alpha_k = \frac{1}{(k+1)^\omega}$, where $\omega \in (0.5, 1)$, the ϵ -optimal performance w.p at least $1 - \delta$ after

$$T = O \left(\left[\frac{\beta^4 R_{max}^2 \log \frac{n\beta R_{max}}{\delta\epsilon}}{\epsilon^2} \right]^{\frac{1}{\omega}} + \left[\beta \log \frac{\beta R_{max}}{\epsilon} \right]^{\frac{1}{1-\omega}} \right)$$

- It has been shown for Q-Learning the finite time PAC bound for a step size, $\alpha_k = \frac{1}{(k+1)^\omega}$, where $\omega \in (0.5, 1)$, the ϵ -optimal performance w.p at least $1 - \delta$ after

$$T = O \left(\left[\frac{\beta^4 R_{\max}^2 \log \frac{n\beta R_{\max}}{\delta\epsilon}}{\epsilon^2} \right]^{\frac{1}{\omega}} + \left[\beta \log \frac{\beta R_{\max}}{\epsilon} \right]^{\frac{1}{1-\omega}} \right)$$

- We see when $\gamma \approx 1$, and we have $O \left(\beta^{\frac{4}{\omega}} + \beta^{\frac{1}{1-\omega}} \right)$ dominates the performance and the optimal bound. Which will be optimized when the exponents are equal, giving us $O \left(\frac{\beta^5}{\epsilon^{2.5}} \right)$

- Summary of Q-Learning vs Speedy Q-Learning

Method	SQL	Q-learning
SC	$\tilde{O}\left(\frac{n\beta^4}{\epsilon^2}\right)$	$\tilde{O}\left(\frac{n\beta^5}{\epsilon^{2.5}}\right)$
CC	$\tilde{O}\left(\frac{n\beta^4}{\epsilon^2}\right)$	$\tilde{O}\left(\frac{n\beta^5}{\epsilon^{2.5}}\right)$
SPC	$\Theta(n)$	$\Theta(n)$

① Motivation for Speedy Q-Learning

② Algorithm

③ Simulations

④ Proposed Modification to SQL

⑤ Results

⑥ Questions

Questions?